

Information Retrieval Project

Part B: Indexing with BERT and Query Interface Development

Group 06

Team Members:

G M Shahariar

Md Taukir Azam Chowdhury

Md. Olid Bhuiyan

Samiha Khan

Zabir Al Nazi

Date of Submission:

March 17th, 2025

Table of Content

Table of Content	2
System Overview	3
Architecture	3
Processing Pipeline:	3
Technologies Used	3
BERT Model & Indexing Schema	3
Query Execution & Ranking	4
Runtime Analysis	4
LLM-based Answers (Bonus)	4
Performance Evaluation with PyLucene	6
Runtime Comparison	6
Ranking Comparison	6
System Limitations	6
Web Interface (Flask-based)	6
Screenshots & Visuals	7
Contribution	8

System Overview

[Github Link](#)

Architecture

The system consists of a dense indexing method:

- **BERT - Transformers (Dense Retrieval):**
 - Transforms text into **contextualized vector representations** for efficient **semantic matching**.
 - Tokenized and truncated Stack Overflow question titles and bodies to first 512 tokens.
 - Used the [CLS] token from BERT's last hidden state for embedding extraction.
- **FAISS (Efficient Similarity Search):**
 - Stores embeddings for **fast nearest neighbor retrieval**
 - **IndexFlatL2** - exact nearest neighbor with euclidean distance.
- **LLM-based Answer:** Using the exacted results, the system generates a natural answer.

Processing Pipeline:

1. **Data Preprocessing:** Clean and tokenize data.
2. **Indexing:** Generate embeddings using BERT and store them in FAISS.
3. **Query Execution:** Retrieve relevant results from FAISS.
4. **Ranking & Output:** Rank and display results.
5. **LLM-based Answer:** Using the exacted results, the system generates a natural answer. (Additional)
6. **Web UI:** The output is displayed using a web UI. (Additional)

Technologies Used

- Web Scraping: Scrapy (for collecting Stack Overflow data).
- Indexing and Search:
 - Sparse Retrieval: Pylucene (Inverted Index, Vector Space Model, Okapi BM25).
 - Dense Retrieval: BERT (🧠 Transformers) and FAISS for semantic search.
- Large Language Model: Qwen 2.5 3B Instruct (via Hugging Face API).
- Backend: Flask API for connecting search and AI models.
- Frontend: HTML, CSS, JavaScript for the user interface.
- Deployment: Hosted on a live server for real-time access.

BERT Model & Indexing Schema

- **Model Choice:** BERT-base-uncased for better accuracy.
- **Passage Splitting:** Each passage is limited to 512 tokens.
- **Embedding Storage:** FAISS index for fast similarity search.

Query Execution & Ranking

- **BERT Querying:** Converts the query into an embedding and retrieves the closest vectors using FAISS.
- **IndexFlatL2** - exact nearest neighbor with euclidean distance.

Runtime Analysis

Step	Time
BERT embedding generation (GPU)	50 minutes
faiss index creation time for bert embeddings	7.32 seconds
Query time	0.101 seconds

More Details:

```
cs242@class-046:~/cs242$ scp znazi002@bolt.cs.ucr.edu:~/bert_faiss_search.py ~/cs242/bert_faiss_search.py
znazi002@bolt.cs.ucr.edu's password:
bert_faiss_search.py
cs242@class-046:~/cs242$ python3 bert_faiss_search.py
  * loading bert tokenizer and model...
  ✓ bert model loaded in 0.4977 sec
  * loading embeddings and metadata...
  ✓ loaded 226340 embeddings with dimension 768.
  ⌚ embedding load time: 0.4108 sec
  ⌚ metadata load time: 6.5972 sec
  * loading existing faiss index...
  ⌚ faiss index load time: 0.7429 sec
  ⌚ total load time: 7.7512 sec

  🔍 computing bert embedding for query: "When to catch java.lang.Error?"...
  ⌚ bert embedding computation time: 0.0774 sec

  🔍 searching for similar questions...

  ⌚ faiss search time: 0.1015 sec
  ⌚ total search time: 0.1017 sec

  * result 1: When to catch java.lang.Error? (distance: 22.4866)
  🔗 link: https://stackoverflow.com/questions/352780/when-to-catch-java-lang-error
  🏷️ tags: ['java', 'error-handling', 'exception', 'java', 'error-handling', 'exception']

  * result 2: When would I use uncaught_exception? (distance: 26.1349)
  🔗 link: https://stackoverflow.com/questions/275249/when-would-i-use-uncaught-exception
  🏷️ tags: ['c++', 'exception', 'error-handling', 'c++', 'exception', 'error-handling']

  * result 3: Connect to BQuery from Python, best practice (distance: 30.4505)
  🔗 link: https://stackoverflow.com/questions/79259176/connect-to-bquery-from-python-best-practice
  🏷️ tags: ['python', 'google-bigquery', 'google-cloud-colab-enterprise', 'python', 'google-bigquery', 'google-cloud-colab-enterprise']
```

Findings:

- BERT is slower but retrieves semantically richer results.

LLM-based Answers (Additional Contribution)

We employ the **Qwen 2.5 3B Instruct**, an instruction-following model. It is hosted on Hugging Face's API inference platform (<https://huggingface.co/Qwen/Qwen2.5-3B-Instruct>), allowing seamless integration with our retrieval pipeline. The model is leveraged to process and generate meaningful answers from the retrieved documents (comments containing solutions). Our intuition is that integrating a Large Language Model (LLM) in the error tag retrieval system will enhance the accuracy and relevance of responses provided to users based on their queries.

To generate relevant responses, the system follows a three-step workflow:

Step 1: Take the user query and documents retrieved by Bert as input

Step 2: The query and the retrieved documents are formatted into a structured prompt. We used the following prompt template -

“““Using the information from the following documents, answer the question concisely:

Document 1: {text_of_retrieved_document_1}

Document 2: {text_of_retrieved_document_2}

Document 3: {text_of_retrieved_document_3}

Question: {user_query}

Answer:

“““

Step 3: The system interacts with the Hugging Face API using the InferenceClient. The model call involves the following steps:

- The prompt is converted to a chat-style conversation template.
- A POST request containing the model name and chat template is sent to the API endpoint.
- The LLM generates a response, leveraging its pre-trained knowledge and the provided context. For generation, the system uses a max token of 1024 tokens.
- The response is streamed in chunks. The system iterates over the output chunks to collect the response gradually. Each chunk is extracted, cleaned (removing None values), and concatenated to form the full response. The final response is presented to the user in a coherent format.

GUI Implementation (Additional Contribution)

Features:

- User-friendly web interface
- Allows Lucene or BERT-based search

- Displays ranked solutions with AI-generated summaries
- Backend: Flask API connecting search & AI models

Hugging Face API Usage:

- Used pre-trained LLM via API call
- Requires API key authentication
- Sends query & retrieved answers for processing

Performance Evaluation with PyLucene

Although PyLucene is not used for indexing in this part, we evaluate the performance of BERT retrieval in comparison to traditional sparse retrieval methods like Lucene.

Runtime Comparison

Method	Indexing Time	Query Time
BERT + FAISS	50 minutes	0.101 seconds
PyLucene	223.26 seconds	0.22 seconds

Ranking Comparison

Example Query	Top-1 BERT Result	Top-1 PyLucene Result
"MySQL group by error"	Logical error discussion	Syntax-related error post

Observation:

- BERT captures intent better for ambiguous queries.
- PyLucene is more precise for well-formed queries with keywords.

System Limitations

- **BERT Limitations:** High memory consumption, slower indexing.
- **PyLucene Limitations:** Does not understand synonyms and semantic similarity.
- **Deployment Challenge:** Hosting FAISS with large datasets requires GPU optimization.

Deployment & Usage Instructions

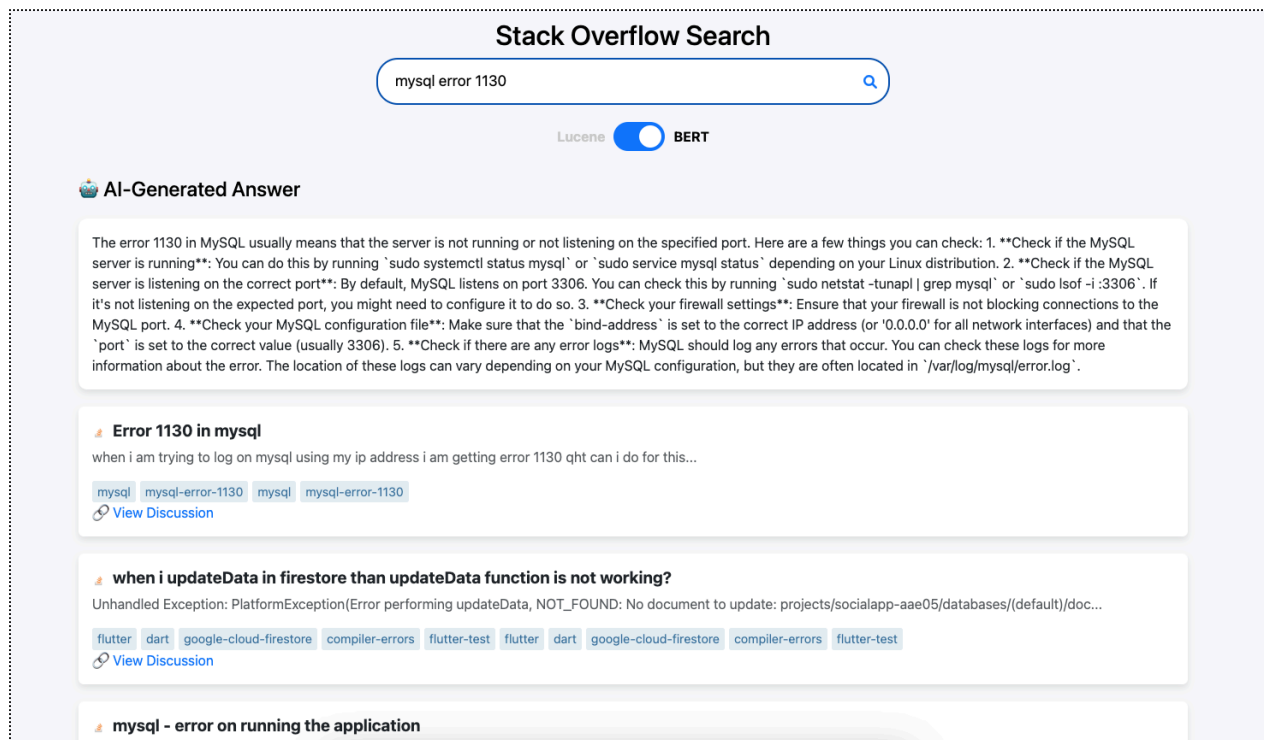
Run the code in terminal: main.py

Web Interface (Flask-based)

1. Start server: `python3 app.py`
2. Access at: <http://169.235.31.51:8080/>

Screenshots & Visuals

Bert-Based Search:



Lucene-Based Search:

Stack Overflow Search

mysql error 1130

Lucene ☒ BERT

AI-Generated Answer

The error 1130 in MySQL usually means that the server is not running or not listening on the specified port. Here are a few things you can check: 1. **Check if the MySQL server is running**: You can do this by running `sudo systemctl status mysql` or `sudo service mysql status` depending on your Linux distribution. 2. **Check if the MySQL server is listening on the correct port**: By default, MySQL listens on port 3306. You can check this by running `sudo netstat -tunapl | grep mysql` or `sudo lsof -i :3306`. If it's not listening on the expected port, you might need to configure it to do so. 3. **Check your firewall settings**: Ensure that your firewall is not blocking connections to the MySQL port. 4. **Check your MySQL configuration file**: Make sure that the `bind-address` is set to the correct IP address (or `'0.0.0.0'` for all network interfaces) and that the `port` is set to the correct value (usually 3306). 5. **Check if there are any error logs**: MySQL should log any errors that occur. You can check these logs for more information about the error. The location of these logs can vary depending on your MySQL configuration, but they are often located in `/var/log/mysql/error.log`.

Error 1130 in mysql

when i am trying to log on mysql using my ip address i am getting error 1130 qht can i do for this ...

mysql-error-1130 mysql

[View Discussion](#)

phpMyAdmin - #1130 - Host 'SERVER' is not allowed to connect to this MySQL server

I run WAMPServer on a Windows server 2012 r2, I want to run a database on it using phpMyAdmin, everything is working but when I try ...

phpmyadmin mysql-error-1130 wampserver

[View Discussion](#)

Output displayed in terminal (With distance score):

```
🔍 computing bert embedding for query: "When to catch java.lang.Error?"...
🕒 bert embedding computation time: 0.0774 sec
🔍 searching for similar questions...
🕒 faiss search time: 0.1015 sec
🕒 total search time: 0.1017 sec

💡 result 1: When to catch java.lang.Error? (distance: 22.4866)
🔗 link: https://stackoverflow.com/questions/352780/when-to-catch-java-lang-error
🏷️ tags: ['java', 'error-handling', 'exception', 'java', 'error-handling', 'exception']

💡 result 2: When would I use uncaught_exception? (distance: 26.1349)
🔗 link: https://stackoverflow.com/questions/275249/when-would-i-use-uncaught-exception
🏷️ tags: ['c++', 'exception', 'error-handling', 'c++', 'exception', 'error-handling']

💡 result 3: Connect to BQuery from Python, best practice (distance: 30.4505)
🔗 link: https://stackoverflow.com/questions/79259176/connect-to-bquery-from-python-best-practice
🏷️ tags: ['python', 'google-bigquery', 'google-cloud-colab-enterprise', 'python', 'google-bigquery', 'google-cloud-colab-enterprise']
```

Github Link:

We have uploaded our full code, with detailed description in the readme file. You can check the code here: [GITHUB](#)

Contribution

BERT Indexing & FAISS Implementation:

- Research and selection of BERT model: Zabir Al Nazi, G M Shahariar
- Implementing BERT-based embedding generation: Md Taukir Azam Chowdhury, Md. Olid Bhuiyan, Samiha Khan
- Storing embeddings in FAISS and optimizing indexing: Zabir Al Nazi
- Improving indexing speed and memory optimization: G M Shahariar

Query Interface Development:

- Command-line interface implementation: Samiha Khan, G M Shahariar
- Web application development (Flask-based): Md Olid Bhuiyan, Md Taukir Azam Chowdhury
- Integration of BERT retrieval with FAISS: Zabir Al Nazi

Performance Analysis & Ranking Comparison:

- Collecting test cases and analyzing query outputs: Zabir Al Nazi
- Comparing runtime efficiency: G M Shahariar

System Deployment & Testing:

- Writing deployment scripts and instructions: Md. Olid Bhuiyan, Samiha Khan, G M Shahariar
- Testing system performance under different query loads: Md Taukir Azam Chowdhury, Zabir Al Nazi, Md. Olid Bhuiyan, Samiha Khan, G M Shahariar

Reporting & Documentation:

- Writing system architecture and indexing details: Md Taukir Azam Chowdhury, Samiha Khan, Zabir Al Nazi
- Explaining query execution and retrieval methods: G M Shahariar, Md Olid Bhuiyan
- Documenting obstacles, solutions, and system limitations: Samiha Khan

LLM-Based Answers (Bonus): G M Shahariar, Md Olid Bhuiyan